



✓ Log odds-ratio

The log odds ratio with an informative (and uninformative) Dirichlet prior (described in [Monroe et al. 2009, Fighting Words](#)) is a common method for finding distinctive terms in two datasets (see [Jurafsky et al. 2014](#) for an example article that uses it to make an empirical argument). This method for finding distinguishing words combines a number of desirable properties:

- it specifies an intuitive metric (the log-odds) for the ratio of two probabilities
- it can incorporate prior information in the form of pseudocounts, which can either act as a smoothing factor (in the uninformative case) or incorporate real information about the expected frequency of words overall.
- it accounts for variability of a frequency estimate by essentially converting the log-odds to a z-score.

In this homework you will implement this ratio for a dataset of your choice to characterize the words that differentiate each one.

```
# Import libraries NLTK and PDF plumber
# download the packages first here
!pip install nltk
!pip install PyPDF2
import nltk

nltk.download('punkt_tab')
import requests # Download from github
from PyPDF2 import PdfReader # PDF reader
from io import BytesIO
import math
import operator
from collections import Counter
```

```
Requirement already satisfied: nltk in /usr/local/lib/python3.12/dist-packages (3.9.1)
Requirement already satisfied: click in /usr/local/lib/python3.12/dist-packages (from nltk) (8.2.1)
Requirement already satisfied: joblib in /usr/local/lib/python3.12/dist-packages (from nltk) (1.5.2)
Requirement already satisfied: regex<=2021.8.3 in /usr/local/lib/python3.12/dist-packages (from nltk) (2024.11.6)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages (from nltk) (4.67.1)
Requirement already satisfied: PyPDF2 in /usr/local/lib/python3.12/dist-packages (3.0.1)
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data] Package punkt_tab is already up-to-date!
```

✓ Part 1

Your first job is to find two datasets with some interesting opposition -- e.g., news articles from CNN vs. FoxNews, books written by Charles Dickens vs. James Joyce, screenplays of dramas vs. comedies. Be creative -- this should be driven by what interests you and should reflect your own originality. **This dataset cannot come from Kaggle.** Feel free to use web scraping (see [here](#) for a great tutorial) or manually copying/pasting text. Aim for more than 10,000 tokens for each dataset.

Save those datasets in two files: "class1_dataset.txt" and "class2_dataset.txt"

Describe each of those datasets and their source in 100-200 words.

```
def pdf_to_txt(url, output_file):
    # Modify the URL to get the raw content from GitHub
    response = requests.get(url)
    reader = PdfReader(BytesIO(response.content))
    text = "\n".join(page.extract_text() or "" for page in reader.pages)
    with open(output_file, "w", encoding="utf-8") as f:
        f.write(text.lower())
```

```
# Download and save the dataset in txt files
pdf_to_txt("https://raw.githubusercontent.com/willie84/anlp-assignments-2025/main/The_Constitution_of_Kenya_2010.pdf", "class1_dataset.txt")
pdf_to_txt("https://raw.githubusercontent.com/willie84/anlp-assignments-2025/main/SACConstitution-web-eng.pdf", "class2_dataset.txt")
```

Type your response here:

The two datasets are the constitutions of Kenya and South Africa. The Kenyan constitution has been obtained as a PDF document from the Kenyan Parliament website here: https://www.parliament.go.ke/sites/default/files/2017-05/The_Constitution_of_Kenya_2010.pdf.

The South African constitution has also been obtained as a PDF from the South African Department of Justice website here: <https://www.justice.gov.za/constitution/SACConstitution-web-eng.pdf>. I have downloaded the PDFs and hosted them on my GitHub to be able to download them and use the PDF library to read the PDF content and convert it to text.

Part 2

Tokenize those texts by filling out the `read_and_tokenize` function below (your choice of tokenizer). The input is a filename and the output should be a list of tokens.

```
def read_and_tokenize(filename: str) -> list[str]:
    """Read the file and output a list of strings (tokens)."""

    from nltk.tokenize import sent_tokenize, word_tokenize

    # Read file
    with open(filename, "r") as f:
        text = f.read()

    # use NLTK word tokenize for tokenizing
    tokens = word_tokenize(text)
    return tokens

# change these file paths to wherever the datasets you created above live.
class1_tokens = read_and_tokenize("class1_dataset.txt")
class2_tokens = read_and_tokenize("class2_dataset.txt")
print(class1_tokens[:10])
print(class2_tokens[:10])
print(len(class1_tokens))
print(len(class2_tokens))

['laws', 'of', 'kenya', 'the', 'constitution', 'of', 'kenya', ',', '2010', 'published']
['the', 'constitution', 'of', 'the', 'republic', 'of', 'south', 'africa', ',', '1996']
61830
59399
```

Part 3

Now let's find the words that characterize each of those sources (with respect to the other). Implement the log-odds ratio with an uninformative Dirichlet prior. This value, $\hat{\zeta}_w^{(i-j)}$ for word w reflecting the difference in usage between corpus i and corpus j , is given by the following equation:

$$\hat{\zeta}_w^{(i-j)} = \frac{\hat{d}_w^{(i-j)}}{\sqrt{\sigma^2(\hat{d}_w^{(i-j)})}}$$

Where:

$$\hat{d}_w^{(i-j)} = \log\left(\frac{y_w^i + \alpha_w}{n^i + \alpha_0 - y_w^i - \alpha_w}\right) - \log\left(\frac{y_w^j + \alpha_w}{n^j + \alpha_0 - y_w^j - \alpha_w}\right)$$

$$\sigma^2(\hat{d}_w^{(i-j)}) \approx \frac{1}{y_w^i + \alpha_w} + \frac{1}{y_w^j + \alpha_w}$$

And:

- y_w^i = count of word w in corpus i (likewise for j)
- $\alpha_w = 0.01$
- V = size of vocabulary (number of distinct word types)
- $\alpha_0 = V * \alpha_w$
- n^i = number of words in corpus i (likewise for j)

In this example, the two corpora are your class1 dataset (e.g., i = your class1) and your class2 dataset (e.g., j = class2). Using this metric, print out the 25 words most strongly aligned with class1, and 25 words most strongly aligned with class2. Again, consult [Monroe et al. 2009, Fighting Words](#) for more detail.

```
def logodds_with_uninformative_prior(tokens_i: list[str], tokens_j: list[str], display=25):
    """Print out the log odds results given two lists of tokens."""
    # Count tokens # Use Counter library
    num_of_tokens_per_word_1 = Counter(tokens_i)
    num_of_tokens_per_word_2 = Counter(tokens_j)

    # Get all types ie unique words in the corpora
    all_types = set(num_of_tokens_per_word_1) | set(num_of_tokens_per_word_2) # Make the union

    # Prior
    Prior_constant = 0.01
```

```

Prior_constant_by_len_of_types = len(all_types) * Prior_constant

dict_of_types_with_log_odds = {}
for typez in all_types:
    type_count_in_corpora_1 = num_of_tokens_per_word_1.get(typez, 0)
    type_count_in_corpora_2 = num_of_tokens_per_word_2.get(typez, 0)

    # As per the class notes
    log_odds = (
        math.log((type_count_in_corpora_1 + Prior_constant) / (len(tokens_i) + Prior_constant_by_len_of_types -
        - math.log((type_count_in_corpora_2 + Prior_constant) / (len(tokens_j) + Prior_constant_by_len_of_types
    )

    # From variance formula

    variance = 1.0 / (type_count_in_corpora_1 + Prior_constant) + 1.0 / (type_count_in_corpora_2 + Prior_constant)

    # Somehow get the Z value
    dict_of_types_with_log_odds[typez] = log_odds / math.sqrt(variance)

# Sort the log odds

sorted_log_odds = sorted(dict_of_types_with_log_odds.items(), key=lambda item: item[1], reverse=True) # descending

# Display results
print("Words that are aligned with document 1 which is Kenyan constitution:")
for typezz, score in sorted_log_odds[:display]:
    print(f"{score:.3f}\t{typezz}")

print("Words that are aligned with document 2 which is South Africa constitution:")
for typezz, score in reversed(sorted_log_odds[-display:]):
    print(f"{score:.3f}\t{typezz}")

```

```
logodds_with_uninformative_prior(class1_tokens,class1_tokens)
```

```

Words that are aligned with document 1 which is Kenyan constitution:
0.000 contents
0.000 ability
0.000 nominate
0.000 thirtieth
0.000 occur
0.000 248
0.000 they
0.000 tabulated
0.000 193
0.000 method
0.000 singly
0.000 none
0.000 murang
0.000 166-appointment
0.000 examination
0.000 two-republic
0.000 substance
0.000 mobility
0.000 211-borrowing
0.000 incurred
0.000 discontinue
0.000 165-high
0.000 nominates
0.000 acquisition
0.000 city165
Words that are aligned with document 2 which is South Africa constitution:
0.000 fee
0.000 margin
0.000 context
0.000 programme
0.000 chairing
0.000 question
0.000 deputy
0.000 won
0.000 before
0.000 reflects
0.000 reflect
0.000 concurrent
0.000 2-independent
0.000 biodiversity
0.000 execute
0.000 entrusted
0.000 includes-
0.000 attainable
0.000 intellectual
0.000 2008
0.000 .192
0.000 representative
0.000 danger

```

```
0.000 146
0.000 exercises
```

To check your work, you can run log-odds on the party platforms from the lab section. With `nlk.word_tokenize` before lower-casing, these should be your top 5 words (and scores, roughly). Depending on your tokenization strategy, your scores might be slightly different.

Democrat:

```
president: 4.75
biden: 4.27
to: 4.11
he: 4.09
has: 4.08
```

Republican

```
republicans: -13.45
our: -11.23
will: -10.88
american: -10.01
restore: -7.97
```

```
!wget --no-check-certificate https://raw.githubusercontent.com/dbamman/anlp25/main/data/2024_democrat_party_platform.txt
!wget --no-check-certificate https://raw.githubusercontent.com/dbamman/anlp25/main/data/2024_republican_party_platform.txt
```

```
--2025-09-09 22:09:29-- https://raw.githubusercontent.com/dbamman/anlp25/main/data/2024_democrat_party_platform.txt
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.109.133, 185.199.110.133, 185.199.111.133
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 283046 (276K) [text/plain]
Saving to: '2024_democrat_party_platform.txt.3'
```

```
2024_democrat_party 100%[=====] 276.41K --.-KB/s in 0.02s
```

```
2025-09-09 22:09:30 (15.2 MB/s) - '2024_democrat_party_platform.txt.3' saved [283046/283046]
```

```
--2025-09-09 22:09:30-- https://raw.githubusercontent.com/dbamman/anlp25/main/data/2024_republican_party_platform.txt
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.111.133, 185.199.110.133, 185.199.111.133
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 35319 (34K) [text/plain]
Saving to: '2024_republican_party_platform.txt.3'
```

```
2024_republican_par 100%[=====] 34.49K --.-KB/s in 0.004s
```

```
2025-09-09 22:09:30 (7.81 MB/s) - '2024_republican_party_platform.txt.3' saved [35319/35319]
```

```
import nltk
logodds_with_uninformative_prior(
    [w.lower() for w in read_and_tokenize("2024_democrat_party_platform.txt")],
    [w.lower() for w in read_and_tokenize("2024_republican_party_platform.txt")]
)
```

Words that are aligned with document 1 which is Kenyan constitution:

```
4.751 president
4.267 biden
4.110 to
4.090 he
4.076 has
3.761 more
3.389 democrats
3.139 for
3.133 also
3.064 administration
3.056 's
2.931 his
2.900 a
2.818 is
2.775 $
2.563 in
2.514 than
2.511 care
2.488 communities
2.297 as
2.275 continue
2.263 working
2.260 work
```

```
2.246  americans
2.235  year
Words that are aligned with document 2 which is South Africa constitution:
-13.446 republicans
-11.235 our
-10.882 will
-10.014 american
-7.966 restore
-7.110 great
-6.325 illegal
-6.072 republican
-5.910 policies
-5.797 :
-5.787 stop
-5.563 again
-5.555 we
-5.551 inflation
-5.506 4
-5.436 must
-5.323 1
-5.304 party
-5.281 3
-5.181 common
-5.154 bring
-5.142 peace
-5.117 5
-5.113 education
-4.944 commitment
```

Start coding or [generate](#) with AI.